Bias and Fairness in Machine Learning: Challenges and Solutions

905F3: WIDER TOPICS IN DATA SCIENCE

Candidate No: 291745

Abstract

Machine learning systems (ML) have started to determine vital choices including recruitment decisions and funding approvals and medical treatments and policing exercises. Such models adopt bias from their historical data sources that produce unfair discriminatory decisions. The origin of ML bias arises from multiple components such as unbalanced data distributions together with system design flaws and structural inequalities which create problems with transparency and fairness in algorithms. The solution to these matters needs a blend of algorithms which perceive fairness along with detection methods for bias as well as supporting legal frameworks. This paper investigates machine learning bias origin points then evaluates both available fairness-oriented research methods and proposed bias reduction solutions. The review of current research demonstrates the significance of explainable AI together with debiasing methods and policy interventions to develop ethical AI systems.

1. Introduction

Processes of information analysis and decision-making have transformed since the invention of machine learning methods. The application of these systems enables better efficiency together with objective outcomes for tasks such as disease diagnosis and loan default predictions and potential job candidate identification. Machine learning systems show no neutrality because they reproduce existed biases in data sets used for training (Barocas et al. 2019). A growing body of studies combines data science with ethical and legal examination of algorithmic fairness because of wide-ranging concerns about the issue.

The origins of bias in ML systems stem from three main factors: structural data inequalities from history and underrepresentation in data samples and the optimization of algorithms for performance above fairness standards (Mehrabi et al., 2021). The establishment of machine learning fairness standards remains unclear to researchers who now ask fundamental questions about this matter. What methods do we have to assess and control problems of biased decision systems? Performance metrics compete with fairness standards which leads to evaluation decisions in machine learning applications.

This essay critically explores both the difficulties and available solutions related to achieving fairness in machine learning applications. The essay starts by identifying bias sources before moving onto definitions and mitigation strategies for fairness and ending with case study analysis and future challenges.

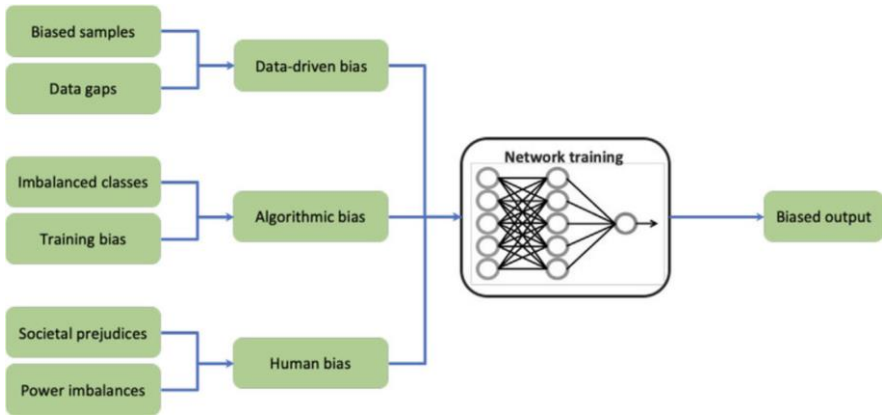

Figure 1: Sources of bias in machine learning systems, including data-driven bias, algorithmic bias, and human bias, each contributing to biased outputs during model training (Norori et al., 2021)

2. The Origins of Bias in Machine Learning

The development pipeline of ML systems presents various opportunities for bias to enter during data collection through deployment. Causes of artificial intelligence bias must be fully understood before creating effective solutions.

2.1 Historical and Societal Bias

Historical bias describes the discriminatory patterns and unequal distribution which exist within historical information. For example, the training of hiring algorithms using historical employment records can result in discrimination against female candidates because men have traditionally dominated specific positions (Bolukbasi et al., 2016). These biases identify societal operational practices that create these patterns which proves challenging to solve exclusively through algorithmic solutions.

2.2 Representation Bias

Models tend to display inferior performance in groups whose numbers remain limited during the training phase. A notable case is that the error rate proved higher for darker-skinned women compared to lighter-skinned male faces in facial recognition systems because training data was inaccurately biased. Research findings demonstrated that model accuracy reached nearly 1.0 for white male and female subjects yet it decreased to roughly 0.75 for black female subjects thus revealing major differences in model performance between population groups (Buolamwini and Gebru, 2018).

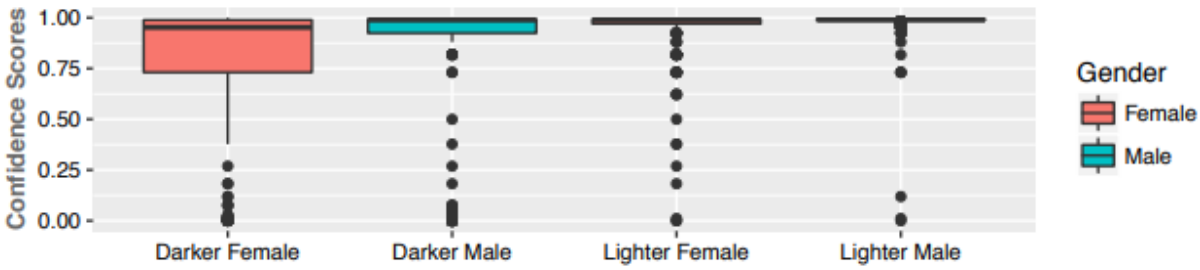

Figure 2: The accuracy of gender classification demonstrates variations between different genders and skin tones across three commercial AI systems (Buolamwini & Gebru, 2018)

2.3 Measurement Bias

A measurement bias develops when variables used to represent actual outcomes contain consistent misrepresentation. A widely used risk prediction algorithm in healthcare used cost of care as a substitute for health needs according to Obermeyer et al. (2019). The utilization of cost of care as a measurement tool resulted in unfair treatment of Black patients because they received less medical care despite having similar illness levels to white patients because of past and systemic inequalities. The healthcare system distributed resources inequitably because the algorithm failed to accurately assess the health requirements of these patients.

2.4 Aggregation Bias and Deployment Bias

The application of population-wide trained models across all subgroups leads to aggregation bias because the models ignore unique characteristics of individual groups. Deployment bias appears because the assumptions from model development fail to match real-world operational conditions.

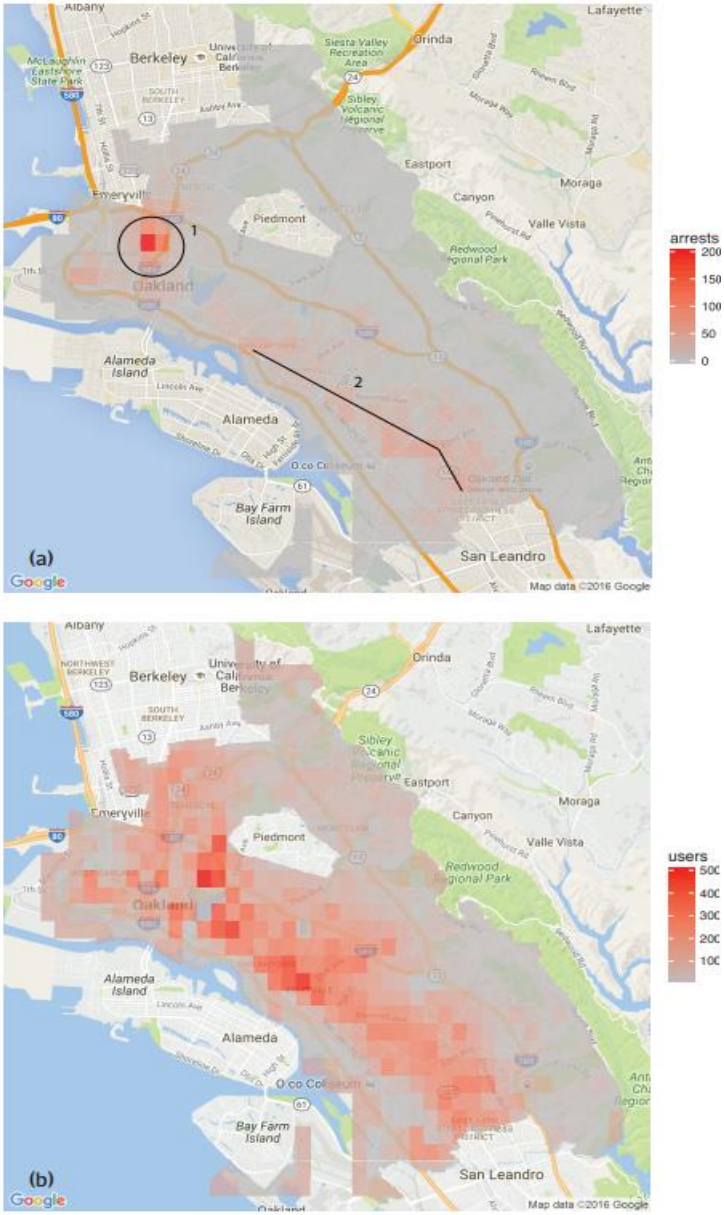

Traditional predictive policing systems demonstrate both types of bias through their reliance on historical crime data that stem from longstanding police over-surveillance focused on low-income neighbourhoods that house mainly racially disenfranchised population groups. The models develop associations between these geographical areas and elevated crime rates. The system deployment leads to higher patrol recommendations for the same communities thus creating additional police contact and documented incidents. Lum and Isaac (2016) show how this feedback system ends up producing biased decisions through the conversion of existing data bias. Figure 3 demonstrates this pattern by showing drug-related arrests against actual drug use locations in Oakland, California. The comparison reveals that police enforcement activities concentrate within particular areas even though the underlying conduct takes place throughout the city.

Figure 3: The relationship between drug arrests (top) against drug use estimates (bottom) throughout Oakland, California (Lum and Isaac, 2016)

3. Defining Fairness in Machine Learning

The definition of fairness in ML remains a complex contested concept. Multiple definitions exist to explain fairness in ML because they represent distinct philosophical and statistical approaches.

3.1 Group Fairness

Group fairness demands equal distribution of results between different demographic groups. The definitions of group fairness consist of multiple formal standards including demographic parity that requires outcomes to be independent from protected characteristics such as race and gender. The second definition under group fairness theory demands that groups receive the same number of true positive outcomes (Hardt et al., 2016). Equalised odds expand upon equal opportunity by demanding equal true and false positive rates within different groups. The definitions exist to prevent one demographic group from receiving persistent unfavourable

treatment. When base rates differ between groups the implementation of one form of group fairness will inevitably lead to the elimination of alternative fairness types (Chouldechova, 2017).

3.2 Individual Fairness

According to Dwork et al. (2012) individual fairness demands that identical individuals should obtain equivalent outcomes. The procedural justice principle of individual fairness faces operational challenges because determining similar individuals remains difficult to define and measure.

3.3 Counterfactual Fairness

The counterfactual fairness evaluation determines if a person would get the same result when their protected characteristic changes while keeping every other factor unchanged (Kusner et al., 2017). The method depends on showing causal relationships through inference to reveal possible biases in hidden structures yet maintains strong requirements about the nature of these relationships.

4. Mitigation Strategies

Research teams have created multiple methods to reduce bias in ML systems. The three main categories of bias reduction techniques include pre-processing methods and in-processing methods and post-processing methods (Feldman et al., 2015).

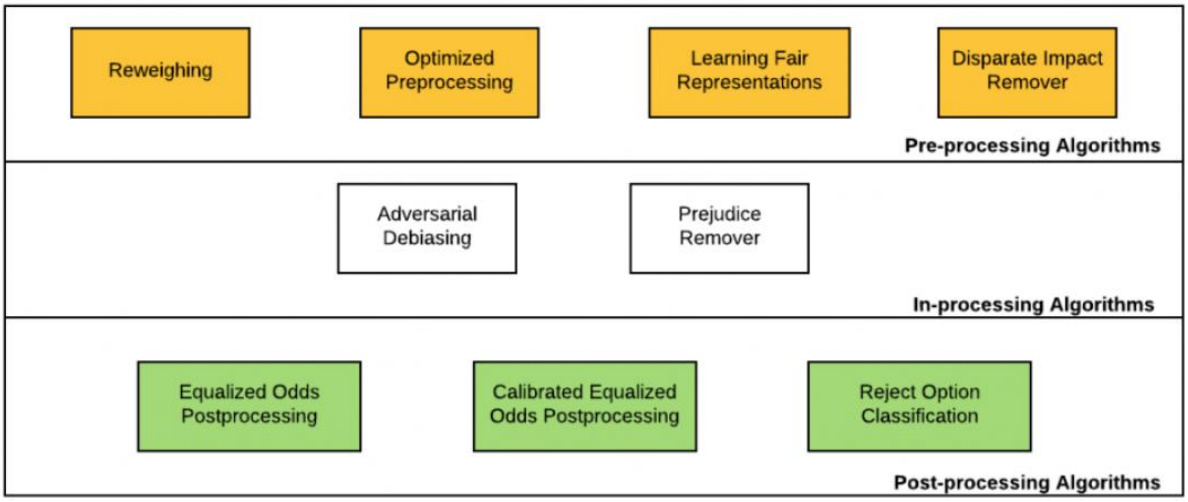

Figure 4: Bias mitigation strategies for machine learning models (Kumar, 2018)

4.1 Pre-processing Methods

The goal of pre-processing techniques is to remove training data bias which should be addressed before building the model. The reweighting technique adjusts training samples to enhance the representation of underrepresented groups according to (Kamiran and Calders, (2012). Fair representation learning constitutes a popular method which transforms data into new spaces that maintain useful information while eliminating sensitive attributes (Zemel et al., 2013).

4.2 In-processing Methods

Model training receives direct fairness constraints through modifications made to the learning algorithm. The adversarial debiasing approach implements a predictive model which serves as an adversary to detect protected attributes. The main model receives optimization from two targets: predictive accuracy and reduced adversary capability to predict sensitive features which minimizes bias (Zhang et al., 2018). The loss function implementation of fairness-constrained optimization embeds equality constraints directly into its calculation.

4.3 Post-processing Methods

The training algorithm gets adjusted to embed fairness constraints during its operational stage. Adversarial debiasing trains a predictive model as an adversary to determine protected attributes. The optimization process for the main model focuses on both accuracy prediction and reducing the adversary's capability to predict sensitive features thus minimizing bias (Zhang et al., 2018). The process of fairness-constrained optimization incorporates specific equality standards directly into the loss function framework.

5. Case Studies

Real-world effects of biased ML systems become evident through multiple widely publicized instances which demonstrate their harmful results.

5.1 COMPAS and Criminal Justice

The COMPAS algorithm which assesses recidivism risk in US courts showed a discriminatory pattern by falsely identifying Black defendants as higher risk than white defendants (Angwin et al., 2016). The incident caused intense public discussion about algorithms hidden methods in judicial systems which showed a competition between accurate predictions and fair treatment.

The evaluation of COMPAS recidivism risk scores through comparison between false positive and false negative rates showed substantial racial disparities. Black defendants faced a double increase in false positive errors where the system incorrectly identified them as high-risk offenders compared to white defendants who experienced more false negative errors showing incorrect low risk assessments (Angwin et al., 2016). Table 1 demonstrates the racial differences in risk level misclassification through its representation of percentage data.

	WHITE	AFRICAN AMERICAN
Labeled Higher Risk, But Didn't Re-Offend	23.5%	44.9%
Labeled Lower Risk, Yet Did Re-Offend	47.7%	28.0%

Table 1: Racial disparities in COMPAS risk misclassification (Angwin et al., 2016)

5.2 Amazon's Recruitment Tool

Amazon eliminated its internal hiring system after it proved to discriminate against female job applicants. The model received training based on past employment records which caused it to reduce resumes containing the term "women's" (Dastin, 2018). The incident demonstrated why researchers need to examine data sets and model results after training.

6. Ethical and Legal Considerations

Fairness through algorithms represents both a technological matter combined with serious moral dilemmas and legal considerations. The General Data Protection Regulation (GDPR) in the European Union gives people the right to understand algorithmic choices through increased transparency (Goodman and Flaxman, 2017). US citizens have some protection against algorithmic discrimination through Fair Credit Reporting Act together with Equal Credit Opportunity Act yet enforcement of these laws remains insufficient. The ethical standards emphasize accountability together with transparency and they require active engagement with stakeholders. According to Floridi et al. (2018) ethical principles should be integrated within the design process of AI systems. The system requires both fair treatment and autonomy preservation as well as harm prevention and human well-being advancement.

7. Challenges and Limitations

Multiple obstacles exist in the path toward achieving fairness in ML systems. Fairness remains in permanent conflict with accuracy levels in ML systems. The implementation of fairness metrics through equal opportunity and demographic parity standards tends to decrease model predictive performance (Corbett-Davies et al., 2017). Developers and stakeholders must choose between effectiveness and equity when this tension arises in their normative decision-making process.

A major drawback exists because researchers lack dependable evidence about what makes an outcome fair. The hiring process and lending activities show that historical data systems contain existing discriminatory patterns. Data that incorporates built-in biases will cause models to recreate those biases regardless of fairness constraint implementation. The process of establishing ethical and situation-dependent fair outcomes represents a substantial ethical and contextual difficulty.

A major obstacle exists regarding interpretability issues. Deep neural networks demonstrate excellent accuracy although they fail to provide any explanations regarding their prediction methods. The untraceable system creates challenges in identifying and fixing unfair patterns particularly when applied in crucial situations including medical care and law enforcement systems. Research into explainable AI techniques continues but these methods do not provide sufficient stability to completely address the identified concerns.

Fairness analysis requires specific context-specific solutions because its interpretation depends on individual conditions. The definition of fairness depends on specific use cases because regional cultural standards as well as legal and social frameworks differ substantially between areas. A fairness intervention designed for one domain could prove either inadequate or dangerous for another domain. Domain experts and stakeholders must actively participate throughout the development cycle of fair ML systems because their experience focuses essential research on system design and evaluation.

8. Future Directions

Current research demonstrates directions which can advance fairness within ML. The field of causal fairness models the variables interrelationships to evaluate hypothetical results. The participatory design frameworks allow affected communities to participate in system development frameworks so their diverse perspectives can shape system behaviour (Veale et

al., 2018). Aequitas and AI Fairness 360 serve as accountability tools that help organizations determine fairness levels in their deployed systems. Effective enforceable standards for algorithmic transparency and fairness need policy and regulatory advancements to be established. This progress hinges on ongoing collaboration among computer scientists, legal professionals, ethicists, and domain-specific experts.

The field shows increasing interest in implementing fairness capabilities directly within development environments that practitioners utilize for their work. Major ML platforms are adding fairness checklists together with built-in metrics dashboards to their systems which guide users to evaluate bias throughout their model development process. Organizational standardization of fairness through operationalization fulfils the essential requirement for transforming ethical intentions into practical procedures. The development of participatory approaches now reaches wider audiences through community-run audit systems that let stakeholders and especially marginalized groups conduct assessments based on their real-world experiences. Academia and industry partners have formed new initiatives to create fairness benchmarks and model cards which promote transparency in their work. These resources provide context for performance results by explaining their boundaries and restrictions which occur when making fairness trade-offs. The upcoming stage in progress will involve comprehensive testing methods that perform fairness measurement through diverse definitions as well as statistical representations of demographic groups. The direction of fair algorithms research has shifted toward practical deployment together with governmental policies as well as sustainable monitoring standards. The evolution of the fairness field requires fundamental changes to ML processes through institutional development which inserts responsible accountability and transparent justice principles from development to deployment.

9. Conclusion

The Logic of Machine Learning create essential problems because these systems now serve as intermediaries to access available resources together with life opportunities. These problems need multiple solutions between technical innovative approaches, ethical thought development, legal protection measures and public participation. Strong efforts toward making ML systems transparent along with being accountable and inclusive will lower the risks and create more trust in these systems despite the unattainable level of perfect fairness. The evolution of technology demands human values to remain central during design processes as well as equal distribution of machine learning benefits. When algorithmic bias receives insufficient attention, it leads to major outcomes that include the breakdown of public trust together with the reinforcement of system-based inequality. The integration of AI in daily decision systems requires fairness to function as a core design objective from the beginning.

References

Lum, K., & Isaac, W. (2016). To predict and serve? *Significance*, 13(5), 14–19.

Angwin, J., Larson, J., Mattu, S., & Kirchner, L. (2016). Machine bias. *ProPublica*.

Barocas, S., Hardt, M., & Narayanan, A. (2019). *Fairness and Machine Learning*. fairmlbook.org.

Bolukbasi, T., Chang, K. W., Zou, J. Y., Saligrama, V., & Kalai, A. (2016). Man is to computer programmer as woman is to homemaker? Debiasing word embeddings. In *Advances in Neural Information Processing Systems*.

Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. In *Proceedings of the 1st Conference on Fairness, Accountability and Transparency*.

Chouldechova, A. (2017). Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2), 153–163.

Corbett-Davies, S., Pierson, E., Feller, A., Goel, S., & Huq, A. (2017). Algorithmic decision making and the cost of fairness. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*.

Dastin, J. (2018). Amazon scraps secret AI recruiting tool that showed bias against women. *Reuters*.

Dwork, C., Hardt, M., Pitassi, T., Reingold, O., & Zemel, R. (2012). Fairness through awareness. In *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*.

Feldman, M., Friedler, S. A., Moeller, J., Scheidegger, C., & Venkatasubramanian, S. (2015). Certifying and removing disparate impact. In *Proceedings of the 21st ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*.

Floridi, L., Cowls, J., Beltrametti, M., Chatila, R., Chazerand, P., Dignum, V., ... & Vayena, E. (2018). AI4People—An ethical framework for a good AI society. *Minds and Machines*, 28(4), 689–707.

Goodman, B., & Flaxman, S. (2017). European Union regulations on algorithmic decision-making and a "right to explanation". *AI Magazine*, 38(3), 50–57.

Hardt, M., Price, E., & Srebro, N. (2016). Equality of opportunity in supervised learning. In *Advances in Neural Information Processing Systems*.

Kamiran, F., & Calders, T. (2012). Data preprocessing techniques for classification without discrimination. *Knowledge and Information Systems*, 33(1), 1–33.

Kusner, M. J., Loftus, J., Russell, C., & Silva, R. (2017). Counterfactual fairness. In *Advances in Neural Information Processing Systems*.

Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., & Galstyan, A. (2021). A survey on bias and fairness in machine learning. *ACM Computing Surveys (CSUR)*, 54(6), 1–35.

Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464), 447–453.

Selbst, A. D., Boyd, D., Friedler, S. A., Venkatasubramanian, S., & Vertesi, J. (2019). Fairness and abstraction in sociotechnical systems. In *Proceedings of the Conference on Fairness, Accountability, and Transparency*.

Veale, M., Van Kleek, M., & Binns, R. (2018). Fairness and accountability design needs for algorithmic support in high-stakes public sector decision-making. In *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*.

Zemel, R., Wu, Y., Swersky, K., Pitassi, T., & Dwork, C. (2013). Learning fair representations. In *International Conference on Machine Learning*.

Zhang, B. H., Lemoine, B., & Mitchell, M. (2018). Mitigating unwanted biases with adversarial learning. In *Proceedings of the 2018 AAAI/ACM Conference on AI, Ethics, and Society*.

Kumar, A. (2018, November 16). *Machine Learning Models – Bias Mitigation Strategies*.

Norori, N., Hu, Q., Aellen, F. M., & Tzovara, A. (2021). *Addressing bias in big data and AI for health care: A call for open science*.